Handoffs a cloud session can actually read

A cloud Claude Code session **clones this repository**. It cannot see anything Git ignores — which includes all of `design-docs/`.

So a handoff written for a cloud session lives **here**, tracked, or it is not reachable at all. Handing a session a `design-docs/handoffs/…` path gives it a path to a file that does not exist in its world.

That mistake has now been made twice: the untracked `CLOUD-1..6` batch on 2026-08-23 and the `CLOUD-7..10` batch on 2026-08-25. Both times all the sessions asked the same question at once, which is the tell.

The rule

- **A handoff for a cloud session goes in `docs/handoffs/`**, is committed, and is pushed before the session is told about it.
- **A handoff for a local session may stay in `design-docs/handoffs/`** — a session on this box reads the disk directly, so the gitignore is irrelevant there.
- Give the session the **in-repo path**, and the one command that reaches it even before the branch is merged:

```
git fetch origin && git show origin/main:docs/handoffs/<file>.md
```

Why these are tracked rather than pasted

Pasting a 300-line prompt into four sessions is slow, error-prone, and painful on a tablet. It also leaves no record of what each session was actually told — and the record is the point: these documents state the acceptance criteria a PR is judged against, and a later reader deserves to see the instructions beside the result.

Historical note, kept deliberately: each of these carries corrections to the plan that produced it, found by checking the repository rather than trusting the issue. Two of the four say the job is not what its issue claims. That is worth reading before writing the next batch.

The 25 August batch — CLOUD-1 … CLOUD-6

Numbered by MERGE ORDER, not by when they were written. "Run 1 first" and "merge 1 first" are deliberately the same number, so a session's name is also its place in the queue.

#	Handoff	Issue(s)	ADR	Merge order
1	`CLOUD-1-issue-495-stash-dtos.md`	#495	0079	**FIRST of all six** — rewrites stash field names everywhere; everything else would pay for its rebase
2	`CLOUD-2-issues-493-494-stash-executor.md`	#493 + #494	0078	after #1

#	Handoff	Issue(s)	ADR	Merge order
3	`CLOUD-3-issue-485-journal-quadratic.md`	#485	0080	**before #6** — shares `activity.rs` and `planner/fetch.rs` with it
4	`CLOUD-4-issue-438-parallel-test-race.md`	#438	0083	independent
5	`CLOUD-5-issues-496-365-fixtures-finished.md`	#496 + #365	0082	independent
6	`CLOUD-6-issue-486-tips-unknown-fold.md`	#486	0081	**after #3** — same two files, and the collision is textual

Only two of the six are truly independent (#4 and #5). The batch's own truth-check found that #3 and #6 both edit `crates/git-vista-core/src/activity.rs` and `crates/git-vista-server/src/planner/fetch.rs` — and not incidentally: they edit *the same doc-comment block*, `activity.rs` ~601-624, which lists two known defects, one each. Each fix deletes its own numbered item and whichever lands second must reword the block header. Both handoffs originally claimed to be independent. Both were wrong.

ADR numbers are assigned here, up front, on purpose. When four sessions run at once and each picks "the next free number", they all pick the same one and the index conflicts. That happened on 25 August. 0074–0077 are taken; this batch claims 0078–0083.

One handoff was written and withdrawn

`CLOUD-4` was originally about **#326** (moving `shape()`'s match arms into their per-operation modules). The truth-check refuted three of its claims — its measurements were eighteen days stale, and issue #326 *explicitly forbids* the single-sweep approach the handoff instructed, for a reason that was live inside this very batch. It is parked at `docs/handoffs/parked/CLOUD-X-issue-326-planner-shape.md`, kept rather than deleted, with the full account in `docs/handoffs/parked/README.md`. #438 took its slot.

A name collision worth knowing about

An **earlier, different** `CLOUD-1 … CLOUD-6` batch was written on 23 August and lives in `design-docs/handoffs/` — untracked, local to Tom's box, and about entirely different issues (#449 capture-refs, cold-build measurement, and so on). Some of those filenames also appear in git history from before `design-docs/` was gitignored.

So `CLOUD-1` names two different jobs depending on which directory you are in. The tracked namespace here — `docs/handoffs/` — is the one that is the record; the numbers restart at 1 per batch, and the batch is identified by this section rather than by an ever-climbing counter. If that ever stops being clear enough, the fix is a dated subdirectory per batch, not a renumber of what has already been handed to a session.

What every handoff in this batch carries

Three things the previous batch proved necessary, in every one of the six:

- **An assigned ADR number**, for the reason above.
- **A stated merge order**, because two of them touch the same files.
- **An explicit instruction to say in the PR body that `ci/browser/run.sh` is unrun** — not to leave it implicit. The browser leg cannot run in a cloud container (`landlock_abi=-1`; the server refuses to start without its strict sandbox tier and INV-13 gives it no degraded mode), and every one of the three defects found in PR #490 was found by that leg and by nothing else.

Never tell a cloud session "`cargo test --workspace` must be green"

It cannot be. Every handoff in the 25 August batch said it, and every one was wrong.

A cloud container fails **320 of 915** `git-vista-server` tests on unmodified `main`, before reaching anything a handoff is about. One cause, printed 535 times in a single run:

> `this operation runs in the strict sandbox tier and this host cannot provide it (missing: landlock_abi>=6, bwrap). Per ADR 0029 the operation is refused rather than run in a weaker tier`

268 of those tests print that refusal verbatim; the other 52 are downstream of it (`CheckFailed { GitSpawnFailed }`, "couldn't run git"). Landlock is absent from that kernel — verified by raw

`landlock_create_ruleset` syscall, independently of the server's own probe — so **installing `bwrap` changes nothing**. `seccomp` and user namespaces are present; the strict tier needs all four.

Five sessions worked around this silently and none of them said so. That is the part that matters: an instruction a session cannot follow does not produce a complaint, it produces a quiet workaround and a PR whose green claim means something different from what the reader assumes.

What to write instead

> Roughly 320 `git-vista-server` tests fail in your container for environmental reasons — the strict sandbox tier is unavailable there. **Run the suite on unmodified `main` first, keep that failing-test set, and compare yours against it. Only the difference is yours.** State the two counts in your PR body. Do not report a sandbox-refusal failure as a defect, and do not claim the suite is green.

The trap that makes this worse than a nuisance

Running the six `CURRENT`-writing tests at 8 threads in a container fails **20/20**. That looks exactly like a reproduction of #438's race, and it is an environment failure — the same set fails at `--test-threads=1`. A session reporting that 20/20 as a repro has reported the wrong bug convincingly.

The session that found this refused to report it, and refused to conclude from `--test-threads=1` failing that the process-global hypothesis was dead, on the grounds that the tests never got far enough to race. Both refusals were right. Its write-up is `CLOUD-4-issue-438-result.md`, and it is the most useful document the batch produced.

The general rule

Before writing an acceptance criterion, ask whether the session you are sending it to can physically satisfy it. A criterion that cannot be met is not a high bar. It is an instruction to improvise, issued to someone with no way to tell you they had to.